

---

# Proj4 – problem1

---



---

과목명	멀티코어컴퓨팅 02분반
교수명	손봉수 교수님
제출일	2023.06.09
학 과	소프트웨어대학 소프트웨어학부
작성자	20200641 임수현

---

(a) environment

OS - window

CPU - 11<sup>th</sup> Gen Intel(R) Core(TM) i7-1195G7

GPU - Colab use.

(b) how to compile

**openmp** : using wsl and g++

```
$g++ -fopenmp openmp_ray.cpp
```

**cuda**: using colab

```
>> !nvcc cuda_ray.cu
```

(c) how to execute

**openmp**: using wsl and g++

```
$/a.out 8 result.ppm
```

**cuda**: using colab

```
>> !./a.out
```

## openmp\_ray.cpp source code

```
1  #include <stdio.h>
2  #include <string.h>
3  #include <stdlib.h>
4  #include <time.h>
5  #include <math.h>
6  #include <iostream>
7  #include <omp.h>
8  #define CUDA 0
9  #define OPENMP 1
10 #define SPHERES 20
11
12 #define rnd( x ) (x * rand() / RAND_MAX)
13 #define INF 2e10f
14 #define DIM 2048
15
16 struct Sphere {
17     float  r,b,g;
18     float  radius;
19     float  x,y,z;
20     float hit( float ox, float oy, float *n ) {
21         float dx = ox - x;
22         float dy = oy - y;
23         if (dx*dx + dy*dy < radius*radius) {
24             float dz = sqrtf( radius*radius - dx*dx - dy*dy );
25             *n = dz / sqrtf( radius * radius );
26             return dz + z;
27         }
28         return -INF;
29     }
30 };
31
32 void kernel(int x, int y, struct Sphere* s, unsigned char* ptr)
33 {
34     int offset = x + y * DIM ;
35     float ox = (x - DIM / 2);
36     float oy = (y - DIM / 2);
37
38     float r = 0, g = 0, b = 0;
39     float maxz = -INF;
40
41     //iterate over spheres
42     for (int i = 0; i < SPHERES; i++) {
43         float n;
44         float t = s[i].hit(ox, oy, &n);
45         if (t > maxz) {
46             float fscale = n;
47             r = s[i].r * fscale;
48             g = s[i].g * fscale;
49             b = s[i].b * fscale;
50             maxz = t;
51         }
52     }
53     //store the color values in the bitmap
54     ptr[offset * 4 + 0] = (int)(r * 255);
55     ptr[offset * 4 + 1] = (int)(g * 255);
56     ptr[offset * 4 + 2] = (int)(b * 255);
57     ptr[offset * 4 + 3] = 255;
58 }
59
```

```

60 void ppm_write(unsigned char* bitmap, int xdim, int ydim, FILE* fp)
61 {
62     int i, x, y;
63     fprintf(fp, "P3\n");
64     fprintf(fp, "%d %d\n", xdim, ydim);
65     fprintf(fp, "255\n");
66     for (y = 0; y < ydim; y++) {
67         for (x = 0; x < xdim; x++) {
68             i = x + y * xdim;
69             fprintf(fp, "%d %d %d ", bitmap[4 * i], bitmap[4 * i + 1], bitmap[4 * i + 2]);
70         }
71         fprintf(fp, "\n");
72     }
73 }
74

```

```

75 int main(int argc, char* argv[])
76 {
77     int no_threads;
78     int option;
79     int x, y;
80     unsigned char* bitmap;
81
82     srand(time(NULL));
83
84     if (argc != 3) {
85         printf("> a.out [option] [filename.ppm]\n");
86         printf("[option] 0: CUDA, 1~16: OpenMP using 1~16 threads\n");
87         printf("for example, '> a.out 8 result.ppm' means executing OpenMP with 8 threads\n");
88         exit(0);
89     }
90     FILE* fp = fopen(argv[2], "w");
91
92     if (strcmp(argv[1], "0") == 0) option = CUDA;
93     else {
94         option = OPENMP;
95         no_threads = atoi(argv[1]);
96     }
97     omp_set_num_threads(no_threads);
98
99     //Allocate memory for spheres
100     struct Sphere* temp_s = (struct Sphere*)malloc(sizeof(struct Sphere) * SPHERES);
101
102     //generate random properties for spheres
103     for (int i = 0; i < SPHERES; i++) {
104         temp_s[i].r = rnd(1.0f);
105         temp_s[i].g = rnd(1.0f);
106         temp_s[i].b = rnd(1.0f);
107         temp_s[i].x = rnd(2000.0f) - 1000;
108         temp_s[i].y = rnd(2000.0f) - 1000;
109         temp_s[i].z = rnd(2000.0f) - 1000;
110         temp_s[i].radius = rnd(200.0f) + 40;
111     }
112
113     //Allocate memory for the bitmap
114     bitmap = (unsigned char*)malloc(sizeof(unsigned char) * DIM * DIM * 4);
115

```

```

116 //start measuring time
117 double start_time= omp_get_wtime();
118
119 //Run OpenMP parallel for loop
120 #pragma omp parallel for private(x,y)
121 for (x = 0; x < DIM; x++){
122     for (y = 0; y < DIM; y++){
123         //call the kernel function
124         kernel(x, y, temp_s, bitmap);
125     }
126 }
127 //stop measuring time
128 double end_time=omp_get_wtime();
129
130 double processing_time = end_time-start_time;
131 std::cout<<"OpenMP ("<<no_threads<<"threads) ray tracing:"<<processing_time<< "sec"<<std::endl;
132 std::cout<<"["<<argv[2]<<"] was generated."<<std::endl;
133
134 //write the bitmap to the ppm file
135 ppm_write(bitmap, DIM, DIM, fp);
136
137 fclose(fp);
138 free(bitmap);
139 free(temp_s);
140
141 return 0;
142 }

```

## cuda\_ray.cu

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>
#include <math.h>

#include <sys/time.h>
#include <cuda.h>
#include <chrono>

using namespace std::chrono;

#define CUDA_0
#define OPENMP_1
#define SPHERES_20
#define DIM_2048

#define rnd(x) (x * rand() / RAND_MAX)
#define INF 2e10f

struct Sphere {
    float r, b, g;
    float radius;
    float x, y, z;
    float (*hit)(struct Sphere* s, float ox, float oy, float* n);
};

// CUDA
__device__ float hit(struct Sphere* s, float ox, float oy, float* n) {
    float dx = ox - s->x;
    float dy = oy - s->y;
    if (dx * dx + dy * dy < s->radius * s->radius) {
        float dz = sqrtf(s->radius * s->radius - dx * dx - dy * dy);
        *n = dz / sqrtf(s->radius * s->radius);
        return dz + s->z;
    }
    return -INF;
}
```

```
// CUDA kernel function
__global__ void kernel(struct Sphere* s, unsigned char* ptr) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    int offset = x + y * DIM;
    float ox = (x - DIM / 2);
    float oy = (y - DIM / 2);

    float r = 0, g = 0, b = 0;
    float maxz = -INF;
    for (int i = 0; i < SPHERES; i++) {
        float n;
        float t = hit(&s[i], ox, oy, &n);
        if (t > maxz) {
            float fscale = n;
            r = s[i].r * fscale;
            g = s[i].g * fscale;
            b = s[i].b * fscale;
            maxz = t;
        }
    }

    ptr[offset * 4 + 0] = (int)(r * 255);
    ptr[offset * 4 + 1] = (int)(g * 255);
    ptr[offset * 4 + 2] = (int)(b * 255);
    ptr[offset * 4 + 3] = 255;
}

void ppm_write(unsigned char* bitmap, int xdim, int ydim, FILE* fp) {
    int i, x, y;
    fprintf(fp, "P3\n");
    fprintf(fp, "%d %d\n", xdim, ydim);
    fprintf(fp, "255\n");
    for (y = 0; y < ydim; y++) {
        for (x = 0; x < xdim; x++) {
            i = x + y * xdim;
            fprintf(fp, "%d %d %d ", bitmap[4 * i], bitmap[4 * i + 1], bitmap[4 * i + 2]);
        }
        fprintf(fp, "\n");
    }
}
```

```

int main(void) {

    unsigned char* bitmap;

    srand(time(NULL));

    FILE* fp = fopen("cudaresult", "w");

    struct Sphere* temp_s = (struct Sphere*)malloc(sizeof(struct Sphere) * SPHERES);
    for (int i = 0; i < SPHERES; i++) {
        temp_s[i].r = rnd(1.0f);
        temp_s[i].g = rnd(1.0f);
        temp_s[i].b = rnd(1.0f);
        temp_s[i].x = rnd(2000.0f) - 1000;
        temp_s[i].y = rnd(2000.0f) - 1000;
        temp_s[i].z = rnd(2000.0f) - 1000;
        temp_s[i].radius = rnd(200.0f) + 40;
    }

    bitmap = (unsigned char*)malloc(sizeof(unsigned char) * DIM * DIM * 4);
    auto start = high_resolution_clock::now();

    //Allocate memory on the GPU
    struct Sphere* dev_s;
    unsigned char* dev_bitmap;
    cudaMalloc((void**)&dev_s, sizeof(struct Sphere) * SPHERES);
    cudaMalloc((void**)&dev_bitmap, sizeof(unsigned char) * DIM * DIM * 4);

    // Copy data from CPU to GPU
    cudaMemcpy(dev_s, temp_s, sizeof(struct Sphere) * SPHERES, cudaMemcpyHostToDevice);

    // Launch the CUDA kernel
    dim3 blocksPerGrid(DIM / 16, DIM / 16);
    dim3 threadsPerBlock(16, 16);
    kernel<<<blocksPerGrid, threadsPerBlock>>>(dev_s, dev_bitmap);

    //Copy the result image data from GPU to CPU
    cudaMemcpy(bitmap, dev_bitmap, sizeof(unsigned char) * DIM * DIM * 4, cudaMemcpyDeviceToHost);

    //Free GPU memory
    cudaFree(dev_s);
    cudaFree(dev_bitmap);

    auto stop = high_resolution_clock::now();
    auto duration = duration_cast<milliseconds>(stop - start);
    printf("CUDA ray tracing :%d sec", (int)duration.count());

    ppm_write(bitmap, DIM, DIM, fp);
    fclose(fp);
    printf("[cudaresult.ppm] was generated.");

    free(bitmap);
    free(temp_s);

    return 0;
}


```




# Test Result

## OPENMP

```
limsoo@limsoo: /mnt/c/Users/limsoo/Desktop/proj4/problem1
limsoo@limsoo:/mnt/c/Users/limsoo/Desktop/proj4/problem1$ g++ -fopenmp openmp_ray.cpp
limsoo@limsoo:/mnt/c/Users/limsoo/Desktop/proj4/problem1$ ./a.out 8 result.ppm
OpenMP (8threads) ray tracing:0.202469sec
[result.ppm] was generated.
```



# thread	1	2	4	6	8
time	460.114	287.943	157.142	139.684	158.794
# thread	10	12	14	16	32
time	143.624	124.994	135.345	120.164	124.558



OpenMP is a programming model for multi-threading, making it easy to implement parallel tasks. Use the `#pragma omp parallel` for statement to run the kernel function as a multi-thread. This allows kernel functions to run simultaneously across multiple threads to parallelize calculations. This

reduces calculation time. Running with OpenMP took 460 msec when calculated with a single thread, but running time was reduced to 120 msec when calculated with a multi-thread using OpenMP. This shows that using OpenMP can improve computation speed through parallel processing.


## CUDA

exec time: 194 msec

```
[11] !nvcc cuda_ray.cu
```

```
!./a.out
```

```
CUDA ray tracing :194 Msec  
[cudaresult.ppm] was generated.
```



CUDA allows you to improve computation speed through parallel processing compared to the CPU. CUDA code leverages the parallel processing power of the GPU to accelerate image rendering.

In the given code, the kernel function runs as the CUDA kernel. The kernel function performs ray tracing for each pixel and calculates the intersection with each sphere to determine the color. Because these calculations are processed in parallel, CUDA allows multiple threads to run simultaneously to speed up the calculation.